

Nonlinear Regression Models

Shovan Biswas

2020/10/30

Libraries

```
library(tidyverse)
library(kableExtra)
library(corrplot)
library(reshape2)
library(Amelia)
library(dlookr)
library(fpp2)
library(plotly)
library(gridExtra)
library(readxl)
library(ggplot2)
library(urca)
library(tseries)
library(AppliedPredictiveModeling)
library(RANN)
library(psych)
library(e1071)
library(corrplot)
library(glmnet)
library(mlbench)
library(caret)
library(earth)
```

Exercise 7.2

Friedman (1991) introduced several benchmark data sets created by simulation. One of these simulations used the following nonlinear equation to create data:

$$y = 10\sin(\pi x_1 x_2) + 20(x_3 - 0.5)^2 + 10x_4 + 5x_5 + N(0, \sigma^2)$$

where the x values are random variables uniformly distributed between $[0, 1]$ (there are also 5 other non-informative variables also created in the simulation). The package **mlbench** contains a function called `mlbench.friedman1` that simulates these data:

(This exercise is based on `library(mlbench)`, which I included in libraries at the top.)

```

set.seed(200)

trainingData <- mlbench.friedman1(200, sd = 1)


## We convert the 'x' data from a matrix to a data frame
## One reason is that this will give the columns names.

trainingData$x <- data.frame(trainingData$x)

## Look at the data using

featurePlot(trainingData$x, trainingData$y)

```



```
## or other methods.
```

```
## This creates a list with a vector 'y' and a matrix  
## of predictors 'x'. Also simulate a large test set to  
## estimate the true error rate with good precision:
```

```
testData <- mlbench.friedman1(5000, sd = 1)  
testData$x <- data.frame(testData$x)
```

Tune several models on these data. For example:

(I included library caret in libraries at the top.)

```
knnModel <- train(x = trainingData$x, y = trainingData$y, method = "knn", preProc = c("center", "scale"),  
knnModel
```

```
## k-Nearest Neighbors  
##  
## 200 samples  
## 10 predictor  
##  
## Pre-processing: centered (10), scaled (10)  
## Resampling: Bootstrapped (25 reps)  
## Summary of sample sizes: 200, 200, 200, 200, 200, 200, ...  
## Resampling results across tuning parameters:  
##  
## k RMSE Rsquared MAE  
## 5 3.466085 0.5121775 2.816838  
## 7 3.349428 0.5452823 2.727410  
## 9 3.264276 0.5785990 2.660026  
## 11 3.214216 0.6024244 2.603767  
## 13 3.196510 0.6176570 2.591935  
## 15 3.184173 0.6305506 2.577482  
## 17 3.183130 0.6425367 2.567787  
## 19 3.198752 0.6483184 2.592683  
## 21 3.188993 0.6611428 2.588787  
## 23 3.200458 0.6638353 2.604529  
##  
## RMSE was used to select the optimal model using the smallest value.  
## The final value used for the model was k = 17.
```

```
knnPred <- predict(knnModel, newdata = testData$x)
```

```
## The function 'postResample' can be used to get the test set  
## performance values  
postResample(pred = knnPred, obs = testData$y)
```

```
## RMSE Rsquared MAE  
## 3.2040595 0.6819919 2.5683461
```

Which models appear to give the best performance? Does MARS select the informative predictors (those named X1–X5)?

Answer:

We observed above that the RMSE of kNN model is 3.2932153. In the following, we'll explore all other models, Neural Networks, MARS and SVM, mentioned in the book, in this order.

Neural Networks

Used code from page 163 of textbook.

```
nnetGrid <- expand.grid(.decay = c(0, 0.01, .1), .size = c(1:10), .bag = FALSE)

nnetTune <- train(x = trainingData$x, y = trainingData$y, method = "avNNnet", tuneGrid = nnetGrid,
                 trControl = trainControl(method = "cv"), preProc = c("center", "scale"),
                 linout = TRUE, trace = FALSE,
                 MaxNWts = 10 * (ncol(trainingData$x) + 1) + 10 + 1,
                 maxit = 500)
```

```
## Warning: executing %dopar% sequentially: no parallel backend registered
```

```
nnetTune
```

```
## Model Averaged Neural Network
##
## 200 samples
## 10 predictor
##
## Pre-processing: centered (10), scaled (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...
## Resampling results across tuning parameters:
##
##  decay  size  RMSE      Rsquared  MAE
##  0.00    1    2.434845  0.7683498  1.921367
##  0.00    2    2.497822  0.7558233  1.993325
##  0.00    3    2.037885  0.8419795  1.609413
##  0.00    4    1.900063  0.8584928  1.529545
##  0.00    5    2.176661  0.8092998  1.628603
##  0.00    6    2.743381  0.7255103  1.988222
##  0.00    7    3.496229  0.6401273  2.493454
##  0.00    8    4.034891  0.5941657  2.749735
##  0.00    9    4.221796  0.5137164  2.800450
##  0.00   10    4.682342  0.5848908  2.818883
##  0.01    1    2.437231  0.7689665  1.934978
##  0.01    2    2.510986  0.7596191  1.988260
##  0.01    3    1.999944  0.8419567  1.555751
##  0.01    4    2.003357  0.8445288  1.549723
##  0.01    5    2.104801  0.8296459  1.664982
##  0.01    6    2.314704  0.7997307  1.857949
##  0.01    7    2.341101  0.8072335  1.872758
##  0.01    8    2.205611  0.8163107  1.748153
##  0.01    9    2.262921  0.8146166  1.776693
##  0.01   10    2.453311  0.7709666  1.981977
```

```
## 0.10 1 2.450897 0.7652309 1.942945
## 0.10 2 2.489399 0.7606443 1.997060
## 0.10 3 2.200693 0.8155496 1.786599
## 0.10 4 2.059322 0.8432340 1.651716
## 0.10 5 2.189025 0.8133603 1.729453
## 0.10 6 2.215091 0.8128993 1.757966
## 0.10 7 2.209521 0.8196474 1.786772
## 0.10 8 2.317124 0.8010433 1.826655
## 0.10 9 2.286711 0.7928430 1.849002
## 0.10 10 2.238560 0.8113030 1.787851
##
```

```
## Tuning parameter 'bag' was held constant at a value of FALSE
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were size = 4, decay = 0 and bag = FALSE.
```

```
Neural_pred <- postResample(pred = predict(nnetTune, newdata = testData$x), obs = testData$y)
Neural_pred
```

```
## RMSE Rsquared MAE
## 2.496722 0.784618 1.685182
```

Observation: RMSE of Neural Networks is 2.496722. It's way higher than what we obtained in kNN (3.2040595).

```
varImp(nnetTune)
```

```
## loess r-squared variable importance
##
## Overall
## X4 100.0000
## X1 95.5047
## X2 89.6186
## X5 45.2170
## X3 29.9330
## X9 6.3299
## X10 5.5182
## X8 3.2527
## X6 0.8884
## X7 0.0000
```

The top 5 variables are X4, X1, X2, X5, X3.

MARS

Used code from page 165 of textbook. Included library(earth) in libraries at the top.

```
marsGrid <- expand.grid(.degree = 1:2, .nprune = 2:38)
marsTuned <- train(x = trainingData$x, y = trainingData$y, method = "earth", tuneGrid = marsGrid,
                  trControl = trainControl(method = "cv"))
marsTuned
```

```

## Multivariate Adaptive Regression Spline
##
## 200 samples
## 10 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...
## Resampling results across tuning parameters:
##
## degree nprune RMSE Rsquared MAE
## 1 2 4.334325 0.2599883 3.607719
## 1 3 3.599334 0.4805557 2.888987
## 1 4 2.637145 0.7290848 2.087677
## 1 5 2.283872 0.7939684 1.817343
## 1 6 2.125875 0.8183677 1.647491
## 1 7 1.766013 0.8733619 1.410328
## 1 8 1.671282 0.8842102 1.324258
## 1 9 1.645406 0.8867947 1.322041
## 1 10 1.597968 0.8926582 1.297518
## 1 11 1.540109 0.8996361 1.237949
## 1 12 1.545349 0.8992979 1.243771
## 1 13 1.535169 0.9010122 1.233571
## 1 14 1.529405 0.9018457 1.223874
## 1 15 1.529405 0.9018457 1.223874
## 1 16 1.529405 0.9018457 1.223874
## 1 17 1.529405 0.9018457 1.223874
## 1 18 1.529405 0.9018457 1.223874
## 1 19 1.529405 0.9018457 1.223874
## 1 20 1.529405 0.9018457 1.223874
## 1 21 1.529405 0.9018457 1.223874
## 1 22 1.529405 0.9018457 1.223874
## 1 23 1.529405 0.9018457 1.223874
## 1 24 1.529405 0.9018457 1.223874
## 1 25 1.529405 0.9018457 1.223874
## 1 26 1.529405 0.9018457 1.223874
## 1 27 1.529405 0.9018457 1.223874
## 1 28 1.529405 0.9018457 1.223874
## 1 29 1.529405 0.9018457 1.223874
## 1 30 1.529405 0.9018457 1.223874
## 1 31 1.529405 0.9018457 1.223874
## 1 32 1.529405 0.9018457 1.223874
## 1 33 1.529405 0.9018457 1.223874
## 1 34 1.529405 0.9018457 1.223874
## 1 35 1.529405 0.9018457 1.223874
## 1 36 1.529405 0.9018457 1.223874
## 1 37 1.529405 0.9018457 1.223874
## 1 38 1.529405 0.9018457 1.223874
## 2 2 4.334325 0.2599883 3.607719
## 2 3 3.599334 0.4805557 2.888987
## 2 4 2.637145 0.7290848 2.087677
## 2 5 2.271844 0.7927888 1.823675
## 2 6 2.114868 0.8200184 1.659485
## 2 7 1.780140 0.8733216 1.429346

```

```
##      2      8      1.663164  0.8891928  1.294968
##      2      9      1.460976  0.9122520  1.180387
##      2     10      1.399692  0.9175376  1.122526
##      2     11      1.380002  0.9216251  1.110556
##      2     12      1.312883  0.9284253  1.063321
##      2     13      1.285612  0.9343029  1.014216
##      2     14      1.328520  0.9286650  1.052185
##      2     15      1.322954  0.9298515  1.045527
##      2     16      1.341454  0.9283961  1.053190
##      2     17      1.344590  0.9280972  1.054209
##      2     18      1.340821  0.9285264  1.050274
##      2     19      1.340821  0.9285264  1.050274
##      2     20      1.340821  0.9285264  1.050274
##      2     21      1.340821  0.9285264  1.050274
##      2     22      1.340821  0.9285264  1.050274
##      2     23      1.340821  0.9285264  1.050274
##      2     24      1.340821  0.9285264  1.050274
##      2     25      1.340821  0.9285264  1.050274
##      2     26      1.340821  0.9285264  1.050274
##      2     27      1.340821  0.9285264  1.050274
##      2     28      1.340821  0.9285264  1.050274
##      2     29      1.340821  0.9285264  1.050274
##      2     30      1.340821  0.9285264  1.050274
##      2     31      1.340821  0.9285264  1.050274
##      2     32      1.340821  0.9285264  1.050274
##      2     33      1.340821  0.9285264  1.050274
##      2     34      1.340821  0.9285264  1.050274
##      2     35      1.340821  0.9285264  1.050274
##      2     36      1.340821  0.9285264  1.050274
##      2     37      1.340821  0.9285264  1.050274
##      2     38      1.340821  0.9285264  1.050274
##
```

```
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nprune = 13 and degree = 2.
```

```
MARS_pred <- postResample(pred = predict(marsTuned, newdata = testData$x), obs = testData$y)
MARS_pred
```

```
##      RMSE  Rsquared      MAE
## 1.2803060 0.9335241 1.0168673
```

Observation: RMSE of MARS is 1.2803060. It's least and best so far.

```
varImp(marsTuned)
```

```
## earth variable importance
##
## Overall
## X1 100.00
## X4 75.33
## X2 48.88
## X5 15.63
## X3 0.00
```

The top 5 variables are X1, X4, X2, X5, X3.

Support Vector Machines

Used code from page 167 of textbook.

```
svmRTuned <- train(x = trainingData$x, y = trainingData$y,
  method = "svmRadial", preProc = c("center", "scale"),
  tuneLength = 14, trControl = trainControl(method = "cv"))

svmRTuned
```

```
## Support Vector Machines with Radial Basis Function Kernel
```

```
##
```

```
## 200 samples
```

```
## 10 predictor
```

```
##
```

```
## Pre-processing: centered (10), scaled (10)
```

```
## Resampling: Cross-Validated (10 fold)
```

```
## Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

##	C	RMSE	Rsquared	MAE
##	0.25	2.504105	0.7940789	1.987142
##	0.50	2.219946	0.8148914	1.750249
##	1.00	2.028115	0.8388693	1.590383
##	2.00	1.899331	0.8561464	1.486326
##	4.00	1.815632	0.8669708	1.424246
##	8.00	1.798299	0.8702910	1.427678
##	16.00	1.797165	0.8702715	1.431259
##	32.00	1.795246	0.8705225	1.429235
##	64.00	1.795246	0.8705225	1.429235
##	128.00	1.795246	0.8705225	1.429235
##	256.00	1.795246	0.8705225	1.429235
##	512.00	1.795246	0.8705225	1.429235
##	1024.00	1.795246	0.8705225	1.429235
##	2048.00	1.795246	0.8705225	1.429235

```
##
```

```
## Tuning parameter 'sigma' was held constant at a value of 0.06104815
```

```
## RMSE was used to select the optimal model using the smallest value.
```

```
## The final values used for the model were sigma = 0.06104815 and C = 32.
```

```
SVM_pred <- postResample(pred = predict(svmRTuned, newdata = testData$x), obs = testData$y)
SVM_pred
```

##	RMSE	Rsquared	MAE
##	2.0693488	0.8263553	1.5718972

Observation: RMSE of SVM is 2.0469184.


```
varImp(svmRTuned)
```

```
## loess r-squared variable importance
##
##      Overall
## X4 100.0000
## X1  95.5047
## X2  89.6186
## X5  45.2170
## X3  29.9330
## X9   6.3299
## X10  5.5182
## X8   3.2527
## X6   0.8884
## X7   0.0000
```

The top 5 variables are X4, X1, X2, X5, X3.

Summary

```
results <- data.frame(t(postResample(pred = knnPred, obs = testData$y))) %>% mutate("Model" = "KNN")
results <- data.frame(t(Neural_pred)) %>% mutate("Model"= "Neural Networks") %>% bind_rows(results)
results <- data.frame(t(MARS_pred)) %>% mutate("Model"= "MARS") %>% bind_rows(results)
results <- data.frame(t(SVM_pred)) %>% mutate("Model"= "Support Vector Machines") %>% bind_rows(results)

results %>% select(Model, RMSE, Rsquared, MAE) %>% arrange(RMSE)
```

```
##           Model      RMSE  Rsquared    MAE
## 1           MARS 1.280306 0.9335241 1.016867
## 2 Support Vector Machines 2.069349 0.8263553 1.571897
## 3      Neural Networks 2.496722 0.7846180 1.685182
## 4              KNN 3.204059 0.6819919 2.568346
```

Conclusion

MARS outperformed the others.

Exercise 7.5

Exercise 6.3 describes data for a chemical manufacturing process. Use the same data imputation, data splitting, and pre-processing steps as before and train several nonlinear regression models.

Answer:

Before we begin answering the questions, let's pre-process the data. In previous Homework 7, this data had to be imputed and Box-coxed. I'll do the same this time, but will not inspect with histograms, because we already know about that.

```
data(CheicalManufacturingProcess)
dim(CheicalManufacturingProcess)
```

```
## [1] 176 58
```

I'll impute the data using the same technique i.e. mice().

```
data_imputed <- mice(CheicalManufacturingProcess, m = 1, method = "pmm", print = F) %>% complete()
```

```
## Warning: Number of logged events: 135
```

```
any(is.na(data_imputed))
```

```
## [1] FALSE
```

At this point, the data is imputed. I'll proceed to Box-Cox it.

```
# Initially identifying columns, whose skewness are not less than 1.
transform_cols <- c()
```

```
for(i in seq(from = 1, to = length(data_imputed), by = 1)) {
  if(abs(skewness(data_imputed[, i])) >= 1) {
    transform_cols <- append(transform_cols, names(data_imputed[i]))
  }
}
```

```
# Applying Box-cox.
```

```
lambda <- NULL
```

```
data_imputed_2 <- data_imputed
```

```
for (i in 1:length(transform_cols)) {
  lambda[transform_cols[i]] <- BoxCox.lambda(abs(data_imputed[, transform_cols[i]]))

  data_imputed_2[c(transform_cols[i])] <- BoxCox(data_imputed[transform_cols[i]], lambda[transform_cols[i]])
}
```

At this point, the data is pre-processed. The pre-processed data is stored in the variable data_imputed_2.

So, I'll proceed to split the data into train and test in 80/20 ratio.

```
set.seed(200)
```

```
split_index <- createDataPartition(data_imputed_2$Yield, p = 0.8, list = FALSE)
```

```
X_train <- data_imputed_2[split_index, ]
```

```
y_train <- data_imputed_2$Yield[split_index]
```

```
X_test <- data_imputed_2[-split_index, ]
```

```
y_test <- data_imputed_2$Yield[-split_index]
```

(a) Which nonlinear regression model gives the optimal resampling and test set performance?

kNN

```
knnModel <- train(x = X_train, y = y_train, method = "knn", preProc = c("center", "scale"), tuneLength = 25)
knnModel
```

```
## k-Nearest Neighbors
##
## 144 samples
## 58 predictor
##
## Pre-processing: centered (58), scaled (58)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 144, 144, 144, 144, 144, 144, ...
## Resampling results across tuning parameters:
##
##  k  RMSE      Rsquared  MAE
##   5  1.210081  0.5601133  0.9444384
##   7  1.200903  0.5694118  0.9409716
##   9  1.203957  0.5686644  0.9411430
##  11  1.220546  0.5574447  0.9575731
##  13  1.234408  0.5520165  0.9643078
##  15  1.258274  0.5344378  0.9900129
##  17  1.266584  0.5344163  1.0004605
##  19  1.272185  0.5356336  1.0082211
##  21  1.279878  0.5371340  1.0153591
##  23  1.295413  0.5283148  1.0270582
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 7.
```

```
knnPred <- postResample(pred = predict(knnModel, newdata = X_test), obs = y_test)
knnPred
```

```
##      RMSE Rsquared      MAE
## 1.3990997 0.6558133 1.1541518
```

Neural Networks

```
set.seed(200)

nnetGrid <- expand.grid(.decay = c(0, 0.01, .1), .size = c(1:10), .bag = FALSE)

nnetTune <- train(x = X_train,
                  y = y_train,
                  method = "avNNet",
                  tuneGrid = nnetGrid,
                  trControl = trainControl(method = "cv"),
```

```

preProc = c("center", "scale"),
linout = TRUE,
trace = FALSE,
MaxNWts = 10 * (ncol(X_train) + 1) + 10 + 1,
maxit = 500)

```

nnetTune

```

## Model Averaged Neural Network
##
## 144 samples
## 58 predictor
##
## Pre-processing: centered (58), scaled (58)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 130, 129, 129, 130, 131, 131, ...
## Resampling results across tuning parameters:
##
##  decay  size  RMSE      Rsquared  MAE
##  0.00    1    1.4751204  0.4111896  1.2105103
##  0.00    2    1.3290201  0.4785784  1.0824154
##  0.00    3    1.4476805  0.6626447  1.0855225
##  0.00    4    1.4011939  0.6021817  1.0737305
##  0.00    5    1.4544828  0.6302205  1.1474825
##  0.00    6    1.8831243  0.4510348  1.5056761
##  0.00    7    2.0811185  0.5146939  1.6128725
##  0.00    8    3.4300327  0.2643533  2.6183578
##  0.00    9    5.1618619  0.3198379  3.4635853
##  0.00   10    5.1964111  0.2722493  3.6164326
##  0.01    1    0.2813606  0.9639782  0.1557607
##  0.01    2    0.4477717  0.9334548  0.2668179
##  0.01    3    0.7257841  0.8620893  0.4595370
##  0.01    4    1.2584315  0.6998765  0.7876371
##  0.01    5    1.2438278  0.7242635  0.8201195
##  0.01    6    0.9605542  0.7539316  0.7604256
##  0.01    7    0.9565450  0.7295212  0.7563796
##  0.01    8    1.1977661  0.6866119  0.8956017
##  0.01    9    1.9605391  0.5404868  1.4096342
##  0.01   10    2.0286259  0.5069970  1.4715894
##  0.10    1    0.4523260  0.9383174  0.3212159
##  0.10    2    1.0942445  0.7856302  0.7113967
##  0.10    3    1.2244591  0.7651362  0.7456411
##  0.10    4    1.1764762  0.7744547  0.7043238
##  0.10    5    1.2866176  0.7286886  0.7297159
##  0.10    6    1.3021583  0.7274787  0.7643164
##  0.10    7    1.1773769  0.7163146  0.8046963
##  0.10    8    1.6074084  0.6065895  1.0162976
##  0.10    9    1.2906389  0.6685646  0.9049381
##  0.10   10    1.1863507  0.6751777  0.9125173
##
## Tuning parameter 'bag' was held constant at a value of FALSE
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were size = 1, decay = 0.01 and bag = FALSE.

```

```
Neural_pred <- postResample(pred = predict(nnetTune, newdata = X_test), obs = y_test)
Neural_pred
```

```
##      RMSE Rsquared      MAE
## 0.6956048 0.8890842 0.2403163
```

MARS

```
marsGrid <- expand.grid(.degree = 1:2, .nprune = 2:38)

marsTuned <- train(x = X_train, y = y_train, method = "earth", tuneGrid = marsGrid, preProc = c("center", "scale"))

marsTuned
```

```
## Multivariate Adaptive Regression Spline
##
## 144 samples
## 58 predictor
##
## Pre-processing: centered (58), scaled (58)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 131, 130, 130, 130, 128, 130, 130, ...
## Resampling results across tuning parameters:
##
## degree nprune RMSE      Rsquared MAE
## 1      2      9.457448e-15 1      9.179015e-15
## 1      3      9.457448e-15 1      9.179015e-15
## 1      4      9.457448e-15 1      9.179015e-15
## 1      5      9.457448e-15 1      9.179015e-15
## 1      6      9.457448e-15 1      9.179015e-15
## 1      7      9.457448e-15 1      9.179015e-15
## 1      8      9.457448e-15 1      9.179015e-15
## 1      9      9.457448e-15 1      9.179015e-15
## 1     10      9.457448e-15 1      9.179015e-15
## 1     11      9.457448e-15 1      9.179015e-15
## 1     12      9.457448e-15 1      9.179015e-15
## 1     13      9.457448e-15 1      9.179015e-15
## 1     14      9.457448e-15 1      9.179015e-15
## 1     15      9.457448e-15 1      9.179015e-15
## 1     16      9.457448e-15 1      9.179015e-15
## 1     17      9.457448e-15 1      9.179015e-15
## 1     18      9.457448e-15 1      9.179015e-15
## 1     19      9.457448e-15 1      9.179015e-15
## 1     20      9.457448e-15 1      9.179015e-15
## 1     21      9.457448e-15 1      9.179015e-15
## 1     22      9.457448e-15 1      9.179015e-15
## 1     23      9.457448e-15 1      9.179015e-15
## 1     24      9.457448e-15 1      9.179015e-15
## 1     25      9.457448e-15 1      9.179015e-15
## 1     26      9.457448e-15 1      9.179015e-15
## 1     27      9.457448e-15 1      9.179015e-15
```

```
## 1      28      9.457448e-15 1      9.179015e-15
## 1      29      9.457448e-15 1      9.179015e-15
## 1      30      9.457448e-15 1      9.179015e-15
## 1      31      9.457448e-15 1      9.179015e-15
## 1      32      9.457448e-15 1      9.179015e-15
## 1      33      9.457448e-15 1      9.179015e-15
## 1      34      9.457448e-15 1      9.179015e-15
## 1      35      9.457448e-15 1      9.179015e-15
## 1      36      9.457448e-15 1      9.179015e-15
## 1      37      9.457448e-15 1      9.179015e-15
## 1      38      9.457448e-15 1      9.179015e-15
## 2        2      9.457448e-15 1      9.179015e-15
## 2        3      9.457448e-15 1      9.179015e-15
## 2        4      9.457448e-15 1      9.179015e-15
## 2        5      9.457448e-15 1      9.179015e-15
## 2        6      9.457448e-15 1      9.179015e-15
## 2        7      9.457448e-15 1      9.179015e-15
## 2        8      9.457448e-15 1      9.179015e-15
## 2        9      9.457448e-15 1      9.179015e-15
## 2       10      9.457448e-15 1      9.179015e-15
## 2       11      9.457448e-15 1      9.179015e-15
## 2       12      9.457448e-15 1      9.179015e-15
## 2       13      9.457448e-15 1      9.179015e-15
## 2       14      9.457448e-15 1      9.179015e-15
## 2       15      9.457448e-15 1      9.179015e-15
## 2       16      9.457448e-15 1      9.179015e-15
## 2       17      9.457448e-15 1      9.179015e-15
## 2       18      9.457448e-15 1      9.179015e-15
## 2       19      9.457448e-15 1      9.179015e-15
## 2       20      9.457448e-15 1      9.179015e-15
## 2       21      9.457448e-15 1      9.179015e-15
## 2       22      9.457448e-15 1      9.179015e-15
## 2       23      9.457448e-15 1      9.179015e-15
## 2       24      9.457448e-15 1      9.179015e-15
## 2       25      9.457448e-15 1      9.179015e-15
## 2       26      9.457448e-15 1      9.179015e-15
## 2       27      9.457448e-15 1      9.179015e-15
## 2       28      9.457448e-15 1      9.179015e-15
## 2       29      9.457448e-15 1      9.179015e-15
## 2       30      9.457448e-15 1      9.179015e-15
## 2       31      9.457448e-15 1      9.179015e-15
## 2       32      9.457448e-15 1      9.179015e-15
## 2       33      9.457448e-15 1      9.179015e-15
## 2       34      9.457448e-15 1      9.179015e-15
## 2       35      9.457448e-15 1      9.179015e-15
## 2       36      9.457448e-15 1      9.179015e-15
## 2       37      9.457448e-15 1      9.179015e-15
## 2       38      9.457448e-15 1      9.179015e-15
##
```

```
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nprune = 2 and degree = 1.
```

```
MARS_pred <- postResample(pred = predict(marsTuned, newdata = X_test), obs = y_test)
MARS_pred
```

```
##          RMSE      Rsquared      MAE
## 7.105427e-15 1.000000e+00 7.105427e-15
```

SVM

```
svmRTuned <- train(x = X_train, y = y_train,
                  method = "svmRadial", preProc = c("center", "scale"),
                  tuneLength = 14, trControl = trainControl(method = "cv"))

svmRTuned
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 144 samples
## 58 predictor
##
## Pre-processing: centered (58), scaled (58)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 131, 128, 128, 131, 129, 129, ...
## Resampling results across tuning parameters:
##
##  C          RMSE      Rsquared  MAE
##  0.25  1.1014916  0.7383513  0.8993541
##  0.50  0.8745268  0.8330189  0.6988335
##  1.00  0.7057515  0.8932710  0.5543114
##  2.00  0.5963458  0.9191256  0.4730832
##  4.00  0.5828162  0.9207911  0.4634073
##  8.00  0.5828068  0.9207979  0.4633994
## 16.00  0.5828068  0.9207979  0.4633994
## 32.00  0.5828068  0.9207979  0.4633994
## 64.00  0.5828068  0.9207979  0.4633994
##128.00  0.5828068  0.9207979  0.4633994
##256.00  0.5828068  0.9207979  0.4633994
##512.00  0.5828068  0.9207979  0.4633994
##1024.00 0.5828068  0.9207979  0.4633994
##2048.00 0.5828068  0.9207979  0.4633994
##
## Tuning parameter 'sigma' was held constant at a value of 0.0122229
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.0122229 and C = 8.
```

```
SVM_pred <- postResample(pred = predict(svmRTuned, newdata = X_test), obs = y_test)
SVM_pred
```

```
##          RMSE  Rsquared      MAE
## 0.7961616 0.9025694 0.5261874
```

Summary

```

results <- data.frame(t(knnPred)) %>% mutate("Model" = "KNN")

results <- data.frame(t(Neural_pred)) %>% mutate("Model"= "Neural Networks") %>% bind_rows(results)

results <- data.frame(t(MARS_pred)) %>% mutate("Model"= "MARS") %>% bind_rows(results)

results <- data.frame(t(SVM_pred)) %>% mutate("Model"= "Support Vector Machines") %>% bind_rows(results)

results %>% select(Model, RMSE, Rsquared, MAE) %>% arrange(RMSE)

```

```

##           Model           RMSE  Rsquared           MAE
## 1           MARS 7.105427e-15 1.0000000 7.105427e-15
## 2 Neural Networks 6.956048e-01 0.8890842 2.403163e-01
## 3 Support Vector Machines 7.961616e-01 0.9025694 5.261874e-01
## 4           KNN 1.399100e+00 0.6558133 1.154152e+00

```

Conclusion

MARS outperformed the others.

- (b) Which predictors are most important in the optimal nonlinear regression model? Do either the biological or process variables dominate the list? How do the top ten important predictors compare to the top ten predictors from the optimal linear model?

```
varImp(marsTuned)
```

```
## earth variable importance
```

```
## Warning in FUN(newX[, i], ...): no non-missing arguments to max; returning -Inf
```

```
##           Overall
## Yield           NaN

```

varImp(marsTuned) did not return any predictor. So, let me try the second best mode i.e. Neural Networks.

```
varImp(nnetTune)
```

```

## loess r-squared variable importance
##
## only 20 most important variables shown (out of 58)
##
##           Overall
## Yield           100.00
## ManufacturingProcess32 38.65
## BiologicalMaterial06 33.94
## ManufacturingProcess13 30.24
## BiologicalMaterial03 29.55
## BiologicalMaterial12 26.73
## ManufacturingProcess17 26.54

```



```
## ManufacturingProcess31    26.05
## ManufacturingProcess09    24.78
## ManufacturingProcess36    24.59
## ManufacturingProcess06    22.09
## BiologicalMaterial02      20.68
## BiologicalMaterial11      18.81
## ManufacturingProcess11    16.95
## ManufacturingProcess29    16.08
## ManufacturingProcess33    15.83
## ManufacturingProcess30    15.17
## BiologicalMaterial09      14.82
## BiologicalMaterial04      14.82
## BiologicalMaterial08      14.26
```

In the case of Neural Networks, 6 of the top ten predictors are ManufacturingProcess predictors and 3 are BiologicalMaterial. So, the ManufacturingProcess predictors dominate.

PLS_MODEL

```
set.seed(200)
pls_model <- train(x = X_train, y = y_train, method = "pls", preProc = c("center", "scale"), trControl = trControl)

pls_prediction <- predict(pls_model, newdata = X_test)

results <- data.frame(Model = "PLS", RMSE = caret::RMSE(pls_prediction, y_test), Rsquared = caret::R2(pls_prediction, y_test))
results
```

```
##   Model      RMSE Rsquared
## 1   PLS 0.3416962 0.9718459
```

```
varImp(pls_model)
```

```
## Warning: package 'pls' was built under R version 3.6.3
```

```
##
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:caret':
##
##   R2
```

```
## The following object is masked from 'package:fpp2':
##
##   gasoline
```

```
## The following object is masked from 'package:corrplot':
##
##   corrplot
```

```
## The following object is masked from 'package:stats':
##
##      loadings

## pls variable importance
##
##      only 20 most important variables shown (out of 58)
##
##              Overall
## Yield              100.00
## ManufacturingProcess32 42.39
## ManufacturingProcess13 35.62
## ManufacturingProcess36 35.22
## ManufacturingProcess17 33.77
## ManufacturingProcess09 32.21
## BiologicalMaterial02   26.84
## ManufacturingProcess12 26.33
## ManufacturingProcess06 26.30
## BiologicalMaterial06   24.60
## BiologicalMaterial08   24.18
## ManufacturingProcess33 23.48
## ManufacturingProcess28 23.46
## BiologicalMaterial12   23.17
## BiologicalMaterial04   22.97
## ManufacturingProcess29 22.31
## ManufacturingProcess31 22.12
## ManufacturingProcess04 21.24
## BiologicalMaterial11   21.12
## BiologicalMaterial03   21.02
```

In nonlinear models, MARS performed best. The RMSE was very close to zero. But PLS_MODEL, which is linear, returned an RMSE of 0.1308365 is higher than MARS's RMSE. So, it faired worse.

However, in linear model PLS_MODEL, among the top 10 variables, ManufacturingProcess is dominant.

- (c) Explore the relationships between the top predictors and the response for the predictors that are unique to the optimal nonlinear regression model. Do these plots reveal intuition about the biological or process predictors and their relationship with yield?


By comparing 20 varImp(nnetTune) with 20 varImp(pls_model), the former being nonlinear and latter liner, we get the following predictors as unique to nonlinear.

BiologicalMaterial03 BiologicalMaterial09 ManufacturingProcess11 ManufacturingProcess30


This was done offline manually on Bash-shell.

The plots of each of these variables that are unique to nonlinear set are shown below.


```
ggplot(X_train, aes(BiologicalMaterial03, Yield)) + geom_point() + ggtitle("BiologicalMaterial03")
```



```
ggplot(X_train, aes(ManufacturingProcess34, Yield)) + geom_point() + ggtitle("BiologicalMaterial09")
```



```
ggplot(X_train, aes(ManufacturingProcess34, Yield)) + geom_point() + ggtitle("ManufacturingProcess11")
```



```
ggplot(X_train, aes(ManufacturingProcess34, Yield)) + geom_point() + ggtitle("ManufacturingProcess30")
```

